

Vehicle Re-Identification for Origin-Destination Flow Estimation:

Bengaluru Mobility Challenge - Phase II

Abhijith Ganesh Prenitha Rajesh Pardheev Krishna

Abstract—This paper presents our approach for vehicle re-identification in the context of the Bengaluru Mobility Challenge (BMC) - Phase II. The objective is to re-identify vehicles seen at one location in another part of the network and estimate origin-destination (O-D) flows over a specific time period. Our solution integrates YOLOv10 object detection, Instance-Batch Normalization (IBN-Net) for feature extraction, and Chroma for efficient vector storage. Using a custom dataset of Indian vehicles, our system achieved competitive performance in re-identification across multiple camera feeds. This document details our methodology, results, and conclusions.

I. INTRODUCTION

The task of vehicle re-identification is essential for estimating origin-destination (O-D) flows, which are crucial for transportation planning and traffic management. The Bengaluru Mobility Challenge (BMC) Phase II evaluates the ability to re-identify vehicles across different locations within a network over a specific time period. Accurate O-D flow estimation is a key factor for effective traffic control and infrastructure development.

In this challenge, participants are tasked with re-identifying vehicles detected by multiple camera feeds, which helps understand traffic patterns, congestion, and vehicle movement between key points. Our approach leverages state-of-the-art detection and feature extraction techniques to enhance the accuracy of re-identification and prediction of O-D flows, im-

proving the overall traffic management system in a dynamic city like Bengaluru.

Through a detailed analysis of the given video footage, we're able to identify the following key points:

- The dataset is substantial, requiring efficient modeling to be incorporated into deep learning networks.
- The data is unstructured and necessitates thorough cleaning and pre-processing.
- Due to the volume of data, it is impractical to load it entirely into memory.
- The data is highly dynamic, demanding real-time updates and processing.

II. METHODOLOGY

Our solution comprises three primary stages: detection, feature extraction, and matching.

A. Detection

We employ the YOLOv10 object detection framework [3] to identify vehicles in video streams. YOLOv10 is known for its real-time detection capability and balance between speed and accuracy.

B. Feature Extraction

For feature extraction, we use the Instance-Batch Normalization Network (IBN-Net) [4], which combines instance and batch normalization to improve the model's generalization capabilities. This enables robust vehicle identification under varying conditions, such as different lighting or weather.

C. Vector Storage with Chroma

Chroma, a vector database optimized for high-dimensional embeddings, is used to store the vehicle feature vectors. This ensures efficient similarity searches between frames captured by different cameras.

D. Matching

We compute cosine similarity between embeddings to match vehicles across multiple frames and locations. A threshold is applied to determine whether the vehicles in different frames are likely to be the same.

III. TRAINING METRICS

We evaluated our solution on a dataset of 24,500 images of Indian vehicles, divided into 19,600 training images and 4,900 testing images. The performance of the models is summarized in Appendix (see Table II).

IV. SOFTWARE AND LIBRARIES

Our project uses the following libraries:

- `json` - for handling data formats.
- `logging` - for logging program events.
- `os` - to interact with the operating system.
- `subprocess` - for running external commands.
- `sys` - to manipulate runtime environment.
- `time` - for time-related functions.
- `chromadb` - for storing feature vectors.
- `cv2` - for image processing.
- `h5py` - for handling large datasets.
- `numpy` - for numerical operations.
- `torch` - for deep learning and neural networks.
- `comet_ml` - for tracking model metrics.
- `PIL` - for image manipulation.
- `torchvision` - for computer vision tasks.

- `ultralytics` - for YOLO object detection.
- **PrenAbhi** - Custom library for the competition guidelines, data handling, and submission.

V. CONCLUSIONS

Our vehicle re-identification system demonstrated strong performance across various model configurations. The combination of YOLOv10 for detection, IBN-Net for feature extraction, and Chroma for vector storage allowed for accurate O-D flow estimation. Future work includes improving detection in challenging environments and optimizing the matching process for larger datasets. The authors of this paper believe that this model is very foundational and needs large-scale Indian/Localized data to improve and detect fine details in the images. At the present stage, the model is able to detect all the commercial classes of vehicles decently enough but needs more data in terms of three-wheelers, bicycles and other non-commercial vehicles.

REFERENCES

- [1] AI City Challenge, <https://www.aicitychallenge.org/>
- [2] Z. Zheng et al., "Connecting Language and Vision for Natural Language-Based Vehicle Retrieval," *arXiv preprint*, 2021. Available at: <https://arxiv.org/pdf/2105.14897>
- [3] A. Wang et al., "YOLOv10: Real-Time End-to-End Object Detection," *arXiv preprint*, 2024.
- [4] X. Pan, P. Luo, J. Shi, and X. Tang, "Two at Once: Enhancing Learning and Generalization Capacities via IBN-Net," in *Proceedings of the European Conference on Computer Vision (ECCV)*, 2018.
- [5] N. Wojke, A. Bewley, and D. Paulus, "Simple Online and Realtime Tracking with a Deep Association Metric," *2017 IEEE International Conference on Image Processing (ICIP)*, pp. 3645-3649, 2017.
- [6] M. Naphade et al., "The 5th AI City Challenge," *2020 IEEE/CVF Conference on Computer Vision and Pattern Recognition Workshops (CVPRW)*, 2020.

- [7] K. He, X. Zhang, S. Ren, and J. Sun, “Deep Residual Learning for Image Recognition,” in Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition (CVPR), 2016, pp. 770-778.

APPENDIX

Here we provide detailed performance metrics for the models and configurations tested.

A. Image Metrics

Below are the images representing various metrics for our model training and evaluation.




Fig. 1: Loss vs Step




Fig. 2: mAP50-95(B) vs Step




Fig. 3: mAP50(B) vs Step




Fig. 4: Precision(B) vs Step




Fig. 5: Recall(B) vs Step




Fig. 6: Validation Box Loss vs Step

TABLE I: Comprehensive Results: Epochs, Loss, Model Parameters, Speed, GFLOPs, and mAP Across All Models

Name	Epochs	Val/Cls_Loss	Model/Parameters	Batch Size	Speed (ms)	Train/Box_Loss	Model/GFLOPs	mAP50(B)
prenag-xl	100	0.748	31,668,362	16	3.907	1.132	171.062	0.944
abhipren-aug	30	0.481	8,071,770	16	1.876	0.878	24.796	0.981
abhipren-noaug	30	0.394	8,071,770	16	1.748	0.613	24.796	0.984
pretrains-nonhyp	25	0.394	8,071,770	28	1.019	0.611	24.796	0.985
pretrains-hyp	28	0.425	8,071,770	-1	0.966	0.681	24.796	0.980

TABLE II: Model Performance Results

Model	mAP50(B)	Speed (ms)	GFLOPs	Val/Cls_Loss
prenag-xl	0.944	3.907	171.062	0.748
abhipren-aug	0.981	1.876	24.796	0.481
abhipren-noaug	0.984	1.748	24.796	0.394
pretrains-nonhyp	0.985	1.019	24.796	0.394
pretrains-hyp ^l	0.980	0.966	24.796	0.425




Fig. 7: Validation Classification Loss vs Step




Fig. 8: Model GFLOPs